



OPEN

Multimodal diffusion framework for collaborative text image audio generation and applications

Junhua Wang¹, Ouya Zhang^{2✉} & Yuan Jiang³

This paper presents a novel framework for collaborative generation across text, image, and audio modalities using an enhanced diffusion model architecture. We introduce a Hierarchical Cross-modal Alignment Network that establishes unified representations while preserving modality-specific characteristics, and a Cross-modal Conditional Diffusion Model that enables flexible generation pathways through innovative conditional embedding and attention-guided mechanisms. Our approach implements cross-modal mutual guidance and consistency optimization to ensure semantic coherence across generated modalities. Experimental evaluations demonstrate significant improvements over state-of-the-art baselines, with an average 11.65% increase in tri-modal semantic alignment. Applications in media content creation, assistive technology, and education show particular promise, with user evaluations confirming enhanced information accessibility and learning experiences. While computational efficiency and domain adaptation remain challenges, this work establishes a foundation for tri-modal collaborative generation that advances multimodal content creation capabilities.

Keywords Multimodal diffusion models, Text-image-audio generation, Cross-modal alignment, Conditional generation, Assistive technology, Media content creation

Research background and development

In recent years, diffusion models have emerged as a revolutionary framework in the field of generative artificial intelligence, demonstrating remarkable capabilities in creating high-quality and diverse content across various modalities¹. The evolution of these models has progressed from single-modality applications to increasingly sophisticated multi-modal approaches, enabling more comprehensive and context-aware content generation². Multi-modal diffusion models represent a significant advancement in the integration of different data types, allowing for the simultaneous processing and generation of text, images, and audio within a unified computational framework³. These models leverage the denoising diffusion probabilistic process to gradually transform random noise into structured content, guided by conditional inputs from various modalities⁴.

The development trajectory of multi-modal diffusion models has been characterized by incremental improvements in architectural design, training methodologies, and cross-modal alignment techniques⁵. Early implementations primarily focused on text-to-image generation, with models such as DALL-E and Stable Diffusion establishing foundational capabilities for transforming textual descriptions into visual representations⁶. Subsequent research expanded these capabilities to include audio generation and cross-modal translations, enabling more comprehensive multi-modal interactions⁷. Recent advancements have focused on enhancing the coherence and semantic consistency across modalities, addressing the challenge of maintaining contextual alignment between textual, visual, and auditory components⁸.

Importance and research challenges

The collaborative generation of text, image, and audio content presents unique opportunities for applications across diverse domains, including entertainment, education, accessibility, and creative industries⁹. Multi-modal content creation enables more immersive and comprehensive information representation, addressing the limitations of single-modality approaches and more closely approximating human perceptual experiences¹⁰. The integration of these modalities can facilitate more intuitive human-computer interactions and enable more expressive communication systems that leverage complementary information channels¹¹.

¹School of Computer Science, South China Business College, Guangdong University of Foreign Studies, Gangzhou 510545, Guangdong, China. ²School of Information Technology and Engineering, Guangzhou College of Commerce, Gangzhou 511363, Guangdong, China. ³School of Marxism, South China Business College, Guangdong University of Foreign Studies, Gangzhou 510545, Guangdong, China. ✉email: zhangouya_gzsyx@126.com

Despite significant progress, several critical challenges persist in multi-modal diffusion research. Cross-modal alignment remains particularly challenging, requiring sophisticated mechanisms to ensure semantic consistency across heterogeneous data types with distinct statistical properties and representational structures¹². Computational efficiency presents another substantial hurdle, as multi-modal generation typically demands greater computational resources compared to single-modality approaches, necessitating optimization strategies for practical applications¹³. Additionally, evaluation methodologies for multi-modal outputs lack standardization, making objective assessment of generation quality and cross-modal coherence difficult¹⁴.

Research objectives and innovations

This research aims to address these challenges by developing an enhanced framework for text-image-audio collaborative generation that prioritizes cross-modal coherence, computational efficiency, and semantic fidelity. The study extends existing approaches^{15,16} by refining attention mechanisms designed for multi-modal contexts and adapting architectural elements that facilitate more effective information exchange between modality-specific components. We build upon recent contrastive diffusion approaches^{16,17} to develop a hierarchical process that enables coordinated generation across modalities while preserving the unique characteristics of each data type.

The main contributions of this work include: (1) an extension of existing cross-modal alignment techniques that integrates contrastive learning within a tri-modal framework, building upon recent advances in bi-modal generation¹⁶; (2) an adaptive sampling strategy that optimizes the diffusion process based on modality-specific quality assessment; and (3) a comprehensive evaluation framework that incorporates both objective metrics and human assessment protocols for multi-modal outputs, with applications in assistive technology and education.

Paper structure

The remainder of this paper is organized as follows: Section II reviews related work in diffusion models and multi-modal generation; Section III presents the proposed methodology and architectural framework; Section IV describes the experimental setup and evaluation protocols; Section V presents and analyzes the results; Section VI discusses limitations and future research directions; and Section VII concludes the paper with a summary of contributions and implications.

Related work and theoretical foundations

Basic principles of diffusion models

Diffusion models represent a class of generative models that have gained significant prominence due to their exceptional generation quality and theoretical elegance¹⁸. The mathematical foundation of these models is rooted in thermodynamics and stochastic calculus, specifically non-equilibrium thermodynamics where a gradual diffusion process transforms a simple distribution into a complex one through a sequence of small perturbations¹⁹. This process is formalized as a Markov chain that systematically adds Gaussian noise to data samples over multiple timesteps.

The forward diffusion process systematically transforms a data sample $\mathbf{x}_0$ into a latent variable $\mathbf{x}_T$ through T steps of noise addition, following the equation:

$$\mathbf{x}_t = \sqrt{\alpha_t} \mathbf{x}_{t-1} + \sqrt{1 - \alpha_t} \boldsymbol{\epsilon}_{t-1}, \boldsymbol{\epsilon}_{t-1} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$$

where $\alpha_t \in (0, 1)$ represents the noise schedule parameter at timestep t , controlling the rate of information destruction²⁰. With a sufficiently large T , the distribution of $\mathbf{x}_T$ approaches an isotropic Gaussian distribution $\mathcal{N}(\mathbf{0}, \mathbf{I})$, effectively erasing all structure from the original data.

The reverse diffusion process, which constitutes the generative component, attempts to recover the original data distribution by learning to invert this noise addition process²¹. This is accomplished by iteratively denoising samples according to:

$$\mathbf{x}_{t-1} = \frac{1}{\sqrt{\alpha_t}} \left(\mathbf{x}_t - \frac{1 - \alpha_t}{\sqrt{1 - \bar{\alpha}_t}} \boldsymbol{\epsilon}_\theta(\mathbf{x}_t, t) \right) + \sigma_t \mathbf{z}$$

where $\mathbf{z} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$, $\bar{\alpha}_t = \prod_{i=1}^t \alpha_i$ and $\boldsymbol{\epsilon}_\theta$ is a neural network trained to predict the noise component²².

The noise prediction mechanism is a crucial element of diffusion models, where a neural network $\boldsymbol{\epsilon}_\theta$ is trained to estimate the noise added during the forward process²³. The training objective function is typically formulated as:

$$\mathcal{L}_{\text{simple}}(\theta) = \mathbb{E}_{t, \mathbf{x}_0, \boldsymbol{\epsilon}} [\|\boldsymbol{\epsilon} - \boldsymbol{\epsilon}_\theta(\mathbf{x}_t, t)\|^2]$$

This objective encourages the model to accurately predict the noise component at each timestep, enabling the gradual recovery of the clean data sample during the reverse diffusion process²⁴.

When compared to Generative Adversarial Networks (GANs), diffusion models offer greater stability during training and typically produce more diverse outputs²⁵. GANs utilize an adversarial training objective where a generator G and discriminator D engage in a minimax game:

$$\min_G \max_D \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} [\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p_{\mathbf{z}}} [\log (1 - D(G(\mathbf{z})))]$$

This formulation often leads to mode collapse and training instability, issues that diffusion models naturally avoid due to their likelihood-based training approach²⁶.

In comparison with Variational Autoencoders (VAEs), diffusion models typically generate higher-quality samples at the cost of increased computational demands²⁷. VAEs optimize a variational lower bound on the data likelihood:

$$\mathcal{L}_{\text{VAE}}(\theta, \phi) = \mathbb{E}_{q_{\phi}(\mathbf{z}|\mathbf{x})} [\log p_{\theta}(\mathbf{x}|\mathbf{z})] - D_{\text{KL}}(q_{\phi}(\mathbf{z}|\mathbf{x}) \parallel p(\mathbf{z}))$$

While VAEs offer efficient sampling, they often produce blurry outputs due to the limitations of their variational approximation²⁸. Diffusion models overcome this limitation by directly modeling the data distribution through a sequence of denoising steps, albeit with increased computational complexity²⁹.

Advances in multimodal learning research

Multimodal learning represents a pivotal research area that aims to process and analyze heterogeneous data from multiple modalities, creating coherent representations that capture cross-modal relationships³⁰. The fundamental challenge in multimodal learning lies in bridging the modality gap—the inherent disparity in statistical properties, dimensions, and semantic structures across different data types such as text, images, and audio³¹. Research in this domain has evolved significantly, with several key representation learning approaches emerging as particularly influential.

Cross-modal mapping techniques establish explicit transformations between different modality spaces, enabling the projection of features from one modality to another³². These approaches typically employ encoder-decoder architectures that learn mappings through paired data supervision:

$$f_{\theta} : \mathcal{X} \rightarrow \mathcal{Y}, \min_{\theta} \mathcal{L}(f_{\theta}(x), y)$$

where f_{θ} represents the mapping function parameterized by θ , transforming data from source modality $\mathcal{X}$ to target modality $\mathcal{Y}$ ³³. While effective for direct translation tasks, these methods often struggle with preserving fine-grained semantic details during transformation, particularly when modalities have substantially different information densities³⁴.

Joint embedding approaches, in contrast, project multiple modalities into a shared semantic space where similar concepts align regardless of their original modality³⁵. These methods typically optimize objective functions that maximize similarity between corresponding cross-modal samples while minimizing similarity between unrelated pairs:

$$\mathcal{L}_{\text{joint}} = \sum_{(x,y) \in \mathcal{P}} d(E_x(x), E_y(y)) - \sum_{(x,y') \in \mathcal{N}} \max(0, m - d(E_x(x), E_y(y')))$$

where E_x and E_y are modality-specific encoders, $d(\cdot, \cdot)$ measures distance in the joint space, $\mathcal{P}$ represents positive pairs, and $\mathcal{N}$ denotes negative pairs³⁶. These approaches have demonstrated impressive performance in retrieval tasks but often lack the generative capabilities needed for high-fidelity cross-modal content creation³⁷.

Attention mechanisms have revolutionized multimodal learning by enabling dynamic, context-dependent interactions between modalities³⁸. The scaled dot-product attention mechanism, foundational to transformer-based architectures, computes:

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$

where Q , K , and V represent queries, keys, and values derived from different modalities or within the same modality³⁹. Cross-attention mechanisms have proven particularly effective for aligning features across modalities, allowing models to focus on relevant cross-modal correspondences when generating or processing multimodal content⁴⁰.

Despite these advances, existing multimodal learning approaches face several limitations in collaborative content generation. First, most approaches exhibit asymmetric performance across different modality transitions, performing well in certain directions (e.g., text-to-image) while struggling with others (e.g., image-to-audio)⁴¹. Second, current methods often encounter modality collapse during training, where the model prioritizes dominant modalities and underutilizes information from others⁴². Third, temporal alignment remains challenging, particularly when synchronizing dynamic modalities like audio with static modalities like images⁴³.

Furthermore, existing approaches typically rely on paired training data, limiting their applicability in scenarios where aligned multimodal datasets are scarce or expensive to obtain⁴⁴. Many current models also struggle with disentangling modality-specific attributes from shared semantic content, complicating tasks that require selective transfer of certain properties across modalities⁴⁵. These limitations highlight the need for more sophisticated frameworks that can facilitate true collaborative generation across text, image, and audio modalities while maintaining semantic consistency and modality-specific fidelity⁴⁶.

Current state of text-image-audio collaborative generation research

Tri-modal collaborative generation involving text, images, and audio represents an emerging frontier in artificial intelligence research, extending beyond traditional dual-modal approaches to create more comprehensive and immersive content⁴⁷. Current technical approaches can be broadly categorized into sequential pipeline

architectures and unified end-to-end frameworks, each presenting distinct advantages and limitations in addressing the inherent complexity of synchronizing three heterogeneous modalities⁴⁸.

Sequential pipeline architectures typically decompose the tri-modal generation process into a series of bi-modal transitions, where outputs from one stage serve as conditional inputs for subsequent stages⁴⁹. These approaches often implement a cascading conditional generation mechanism that can be formalized as:

$$p(y_{\text{text}}, y_{\text{image}}, y_{\text{audio}} | c) = p(y_{\text{text}} | c) \cdot p(y_{\text{image}} | y_{\text{text}}, c) \cdot p(y_{\text{audio}} | y_{\text{text}}, y_{\text{image}}, c)$$

where c represents the initial condition, and each modality is generated conditionally based on previously generated modalities⁵⁰. While this approach simplifies the generation process, it suffers from error propagation, where inaccuracies in early stages compound throughout the pipeline, potentially resulting in semantic inconsistencies across modalities⁵¹.

Unified end-to-end frameworks, conversely, attempt to model the joint distribution of all three modalities simultaneously through shared latent representations or cross-modal attention mechanisms⁵². These approaches typically employ a joint optimization objective:

$$\mathcal{L}_{\text{joint}} = \lambda_{\text{text}} \mathcal{L}_{\text{text}} + \lambda_{\text{image}} \mathcal{L}_{\text{image}} + \lambda_{\text{audio}} \mathcal{L}_{\text{audio}} + \lambda_{\text{align}} \mathcal{L}_{\text{alignment}}$$

where modality-specific losses ($\mathcal{L}_{\text{text}}$, $\mathcal{L}_{\text{image}}$, $\mathcal{L}_{\text{audio}}$) are combined with an alignment loss ($\mathcal{L}_{\text{alignment}}$) weighted by respective coefficients⁵³. While conceptually more elegant, these approaches face significant computational challenges and often require complex architectural designs to effectively model cross-modal dependencies.

Maintaining semantic consistency across three modalities presents a particularly formidable challenge in tri-modal generation systems⁵⁴. Current approaches employ various techniques to address this challenge, including cross-modal attention mechanisms, contrastive learning objectives, and semantic bridging networks that explicitly model relationships between modality-specific features⁵⁵. Recent research has explored consistency regularization terms that enforce alignment between modalities:

$$\mathcal{L}_{\text{consistency}} = \mathbb{E}_{x,y,z \sim p_{\text{data}}} [d(f_{\text{text}}(x), f_{\text{image}}(y), f_{\text{audio}}(z))]$$

where f_{modality} represents modality-specific encoders and $d(\cdot, \cdot, \cdot)$ measures tri-modal alignment⁵⁶. However, defining appropriate metrics for tri-modal alignment remains an open research question, with existing approaches often resorting to pairwise alignment metrics that may not fully capture the complex interactions among all three modalities.

Quality assessment methodologies for tri-modal generated content present another significant challenge, with existing evaluation frameworks largely adapted from bi-modal contexts and potentially inadequate for capturing the full spectrum of tri-modal interactions⁵⁷. Current evaluation approaches typically employ a combination of modality-specific metrics (e.g., FID for images, BLEU for text) alongside cross-modal alignment metrics and human evaluations⁵⁸. However, these methods often fail to holistically assess the coherence and semantic consistency across all three modalities simultaneously.

Several critical limitations persist in current tri-modal generation research. First, most existing approaches exhibit modal bias, performing substantially better on dominant modalities while struggling with others, particularly audio generation⁵⁹. Second, temporal alignment between static (text, images) and dynamic (audio) modalities remains poorly addressed, with current solutions often producing temporally inconsistent outputs⁶⁰. Third, current frameworks typically require large volumes of aligned tri-modal data, which remains scarce compared to bi-modal datasets⁶¹.

The present research addresses these limitations by proposing a novel hierarchical diffusion framework that explicitly models both modality-specific characteristics and cross-modal dependencies, introducing an adaptive attention mechanism that dynamically balances information flow across modalities, and developing a comprehensive evaluation methodology specifically designed for tri-modal outputs⁶². By addressing these gaps, our work aims to advance the state-of-the-art in collaborative generation across text, image, and audio modalities.

Multimodal diffusion model collaborative generation mechanism

Tri-modal joint representation learning framework

This section presents a novel framework for tri-modal joint representation learning that addresses the fundamental challenges of unifying text, image, and audio representations within a shared semantic space. The proposed Hierarchical Cross-modal Alignment Network (HiCAN) employs a multi-level encoding strategy that progressively aligns features across modalities while preserving modality-specific characteristics essential for high-fidelity generation. Unlike previous approaches that primarily focus on pairwise modality alignment, HiCAN explicitly models the complex interactions among all three modalities simultaneously.

Figure 1 illustrates the overall architecture of our proposed framework, showing the data flow from modality-specific encoders through the alignment network to the unified representation space. The multi-level attention mechanism enables dynamic information exchange between modalities at varying levels of abstraction.

The architecture consists of three primary components: (1) modality-specific encoders that extract rich features from each input type, (2) a cross-modal alignment module that projects these features into a unified representation space, and (3) a hierarchical semantic fusion mechanism that captures complex relationships across modalities at multiple levels of abstraction. This design enables the framework to maintain a balance between modality-specific feature preservation and cross-modal semantic alignment.




Fig. 1. Hierarchical Cross-modal Alignment Network (HiCAN) Architecture. The framework processes text, image, and audio inputs through modality-specific encoders (left), projects features to a unified space through contrastive alignment (middle), and performs cross-modal attention at multiple levels to generate the final unified representation (right).

For initial feature extraction, each modality is processed through specialized encoders optimized for the corresponding data type. The text modality employs a transformer-based encoder that captures contextual relationships between words and phrases. Image content is processed through a vision transformer architecture that extracts both local and global visual features. Audio signals are encoded using a temporal convolutional network that captures frequency patterns and temporal dynamics. Formally, the encoding process for each modality can be expressed as:

$$\begin{aligned} E_{\text{text}}(x_{\text{text}}) &= z_{\text{text}} \in \mathbb{R}^{d_{\text{text}} \times l_{\text{text}}} \\ E_{\text{image}}(x_{\text{image}}) &= z_{\text{image}} \in \mathbb{R}^{d_{\text{image}} \times h \times w} \\ E_{\text{audio}}(x_{\text{audio}}) &= z_{\text{audio}} \in \mathbb{R}^{d_{\text{audio}} \times t} \end{aligned}$$

where x_{modality} represents the input data for each modality, z_{modality} denotes the extracted features, and dimensions d_{modality} , l_{text} , h , w , and t represent feature dimensionality and spatial/temporal dimensions respectively.

The cross-modal alignment module then projects these heterogeneous features into a unified representation space through a series of modality-specific transformation networks. Rather than forcing direct alignment, which can lead to information loss, the proposed approach employs a contrastive alignment mechanism that preserves modality-specific granularity while ensuring semantic correspondence. The transformation networks adapt each modality's feature space to a common dimensionality d_{unified} while maintaining the original structural properties:

$$\Phi_{\text{modality}} : \mathbb{R}^{d_{\text{modality}} \times *} \rightarrow \mathbb{R}^{d_{\text{unified}} \times *}$$

where $*$ represents the varying dimensional structures of each modality. These transformations are optimized to preserve both intra-modal relationships and cross-modal semantic correspondences through a temperature-scaled InfoNCE loss:

$$\mathcal{L}_{\text{contrastive}} = -\log \frac{\exp(\text{sim}(z_i, z_j^+) / \tau)}{\exp(\text{sim}(z_i, z_j^+) / \tau) + \sum_{k=1}^K \exp(\text{sim}(z_i, z_k^-) / \tau)}$$

where z_i and z_j^+ are embeddings from different modalities corresponding to the same semantic content (positive pairs), z_k^- represents embeddings from different content (negative samples), $\text{sim}(\cdot, \cdot)$ is cosine similarity, and τ is a temperature parameter (set to 0.07). Negative samples are obtained through both in-batch negatives and a memory bank of 8,192 features updated with a momentum of 0.999. The contrastive loss is combined with the diffusion objective through $\mathcal{L}_{\text{total}} = \mathcal{L}_{\text{diffusion}} + \lambda_{\text{contrastive}} \mathcal{L}_{\text{contrastive}}$ where $\lambda_{\text{contrastive}} = 0.5$ balances the two objectives.

The hierarchical semantic fusion mechanism forms the core innovation of the proposed framework, implementing a multi-level attention architecture that captures cross-modal relationships at varying levels of abstraction. At each level l in the hierarchy, the framework computes a tri-modal attention tensor A^l that explicitly models the interactions between all three modalities:

$$A^l = \text{softmax} \left(\frac{Q_{\text{tri}}^l (K_{\text{tri}}^l)^T}{\sqrt{d_k}} \right)$$

where Q_{tri}^l and K_{tri}^l represent the query and key matrices derived from the concatenated features of all three modalities at level l . This attention mechanism enables the model to capture complex dependencies that may exist across specific elements of the three modalities simultaneously.

The final unified representation is constructed through a progressive fusion process that combines information across modalities while respecting their hierarchical relationships. This is achieved through a weighted combination of modality-specific features and cross-modal attention outputs:

$$z_{\text{unified}} = \sum_{l=1}^L \alpha_l [\beta_{\text{text}}^l z_{\text{text}}^l + \beta_{\text{image}}^l z_{\text{image}}^l + \beta_{\text{audio}}^l z_{\text{audio}}^l + \gamma^l A^l V_{\text{tri}}^l]$$

where α_l represents the importance weight for hierarchical level l , $\beta_{\text{modality}}^l$ controls the contribution of each modality at level l , γ^l weights the cross-modal attention information, and V_{tri}^l is the value matrix for the tri-modal attention computation. These importance weights are learned jointly through a dedicated lightweight prediction network that outputs normalized weights via softmax. Table 1 presents an ablation study on different modality weight configurations, demonstrating the effectiveness of our learned weighting approach compared to fixed configurations.

As shown in Table 1, learned weights consistently outperform fixed configurations across all evaluation metrics.

The proposed framework incorporates a semantic consistency regularization term during training that explicitly encourages alignment between corresponding semantic concepts across modalities. This regularization is formulated as:

$$\begin{aligned} \mathcal{L}_{\text{consistency}} = & \sum_{(x,y,z) \in \mathcal{D}} [f_{\text{text}}(x) - f_{\text{image}}(y)]^2 + [f_{\text{text}}(x) - f_{\text{audio}}(z)]^2 \\ & + [f_{\text{image}}(y) - f_{\text{audio}}(z)]^2 \end{aligned}$$

where f_{modality} are projection functions mapping each modality to the shared semantic space, and (x, y, z) represents aligned triplets from the dataset $\mathcal{D}$. This regularization complements the contrastive objective by explicitly minimizing distances between aligned modalities, ensuring the unified representation preserves semantic coherence while maintaining modality-specific characteristics necessary for high-quality generation. By addressing both representation unification and semantic relationship modeling in an integrated framework, HiCAN provides a solid foundation for the subsequent diffusion-based generation processes described in the following sections.

Cross-modal conditional diffusion model design

Building upon the tri-modal joint representation learning framework, we propose a novel Cross-modal Conditional Diffusion Model (CCDM) that enables coherent generation across text, image, and audio domains. The unified representation z_{unified} produced by the HiCAN framework directly serves as the conditional embedding c in our diffusion model, establishing a seamless connection between representation learning and the generative process. CCDM extends traditional diffusion models by incorporating multi-modal conditional controls that guide the generative process toward semantically consistent outputs across all three modalities. Table 2 presents the detailed architectural components of the proposed model.

The conditional embedding mechanism forms the core innovation of CCDM, enabling effective incorporation of information from source modalities into the target generation process. Rather than simply concatenating or adding conditional information to the diffusion input, we introduce a hierarchical conditioning approach that

Weight configuration	TSAS	FID	FAD	BLEU-4
Equal weights (0.33)	0.689	21.4	2.73	0.237
Text-biased (0.5, 0.25, 0.25)	0.723	23.1	2.81	0.251
Image-biased (0.25, 0.5, 0.25)	0.701	18.6	2.95	0.226
Audio-biased (0.25, 0.25, 0.5)	0.677	24.3	2.42	0.214
Learned weights (ours)	0.782	17.9	2.38	0.263

Table 1. Ablation study on modality weight parameters (β values).

Module name	Input dimensions	Output dimensions	Parameters (millions)	Functional description
Text encoder	$1 \times 512 \times 768$	$1 \times 77 \times 1024$	123.7	Transformer-based encoder mapping text tokens to high-dimensional embeddings with contextual understanding
Image encoder	$3 \times 256 \times 256$	$1 \times 256 \times 1024$	86.4	Vision transformer converting images to patch-based representations with spatial relationships preserved
Audio encoder	$1 \times 16,000 \times 1$	$1 \times 512 \times 1024$	42.3	Temporal CNN with self-attention layers extracting frequency and temporal features from audio signals
Joint feature mapping layer	$(77 + 256 + 512) \times 1024$	845×2048	14.8	Projects heterogeneous features to unified representation space while preserving modal-specific characteristics
Condition injection module	845×2048	$16 \times 16 \times 1024$	9.2	Transforms unified embeddings into spatial-temporal conditional signals for diffusion process control
Diffusion U-Net	$16 \times 16 \times 1024$	Varies by modality	768.5	Hierarchical denoising network with cross-modal attention at multiple scales supporting conditional generation
Sampling controller	Varies by modality	Varies by modality	2.1	Adaptive sampling mechanism adjusting noise schedules based on generation quality metrics for each modality

Table 2. Cross-modal conditional diffusion model structure. Total parameters: 1,047.0 million.

injects modal information at multiple levels of the denoising network. The conditional diffusion process can be formalized as:

$$p_{\theta}(x_{0:T}|c) = p(x_T) \prod_{t=1}^T p_{\theta}(x_{t-1}|x_t, c)$$

where c represents the conditional information derived from the source modalities. The transition probability $p_{\theta}(x_{t-1}|x_t, c)$ is modeled as a Gaussian distribution:

$$p_{\theta}(x_{t-1}|x_t, c) = \mathcal{N}(x_{t-1}; \mu_{\theta}(x_t, t, c), \Sigma_{\theta}(x_t, t, c))$$

The conditional mean μ_{θ} and variance Σ_{θ} are predicted by the diffusion network with explicit dependence on the conditional information c .

The attention-guided diffusion process extends standard diffusion by incorporating cross-modal attention blocks throughout the U-Net architecture. These blocks enable the model to selectively focus on relevant conditional information during the denoising process. The cross-modal attention mechanism at layer l is defined as:

$$\text{CrossAttn}^l(h^l, c) = \text{SoftMax}\left(\frac{Q^l(h^l) \cdot K^l(c)^T}{\sqrt{d_k}}\right) \cdot V^l(c)$$

where h^l represents the hidden features at layer l of the diffusion U-Net, c is the conditional information, and Q^l , K^l , and V^l are learnable projections. This mechanism allows the model to dynamically attend to different aspects of the conditional information based on the current generation state.

To enhance generation fidelity, we introduce a novel quality-adaptive adjustment method that dynamically modifies the sampling strategy based on real-time quality assessment. The adaptive sampling process implements a modality-specific noise schedule:

$$\alpha_t^m = \alpha_{\text{base}} \cdot (1 + \delta^m \cdot \mathcal{Q}^m(x_t))$$

where α_t^m is the adjusted step size for modality m at timestep t , α_{base} is the base step size, δ^m is a modality-specific adjustment factor, and $\mathcal{Q}^m(x_t)$ is a quality assessment function that evaluates the current generation quality. This approach allows the model to allocate more computational resources to challenging aspects of the generation process.

The integration of multiple conditional modalities is achieved through a weighted fusion mechanism that balances the influence of each source modality on the target generation. The noise prediction network incorporates this fused conditional information as:

$$\epsilon_{\theta}(x_t, t, c) = \epsilon_{\theta}\left(x_t, t, \sum_{m \in M_{\text{source}}} w_m \cdot \Phi_m(c^m)\right)$$

where M_{source} represents the set of source modalities, w_m is the importance weight for modality m , c^m is the conditional information from modality m , and Φ_m is a modality-specific transformation function that adapts the conditional information to the appropriate format.

The proposed CCDM architecture supports flexible generation pathways, enabling various conditional generation scenarios including text-to-image-audio, image-to-text-audio, and audio-to-text-image. This flexibility is achieved through the modular design of the condition injection module and the adaptive sampling

controller, which can be configured to prioritize different modalities based on the specific generation task. The model maintains semantic consistency across modalities through the joint feature mapping layer that establishes explicit correspondences between concepts in different modality spaces.

Multi-modal collaborative generation algorithm

Leveraging the cross-modal conditional diffusion model architecture described in the previous section, we develop a comprehensive Multi-modal Collaborative Generation (MCG) algorithm that enables flexible generation pathways across text, image, and audio modalities. The proposed algorithm supports arbitrary conditional generation scenarios, allowing any subset of modalities to serve as conditioning inputs for generating the remaining modalities. Table 3 provides a detailed overview of the algorithm’s step-by-step process.

The core of our approach lies in the unified conditional generation mechanism that enables the synthesis of any target modality based on arbitrary source modalities. Unlike previous approaches that typically implement fixed generation pathways (e.g., text-to-image or text-to-audio), our algorithm dynamically constructs the generation process based on the available source modalities. The conditional generation process for a target modality j given source modalities M_{src} can be formalized as:

$$p_{\theta} \left(x_{0:T}^j | \{ x_i^{src} \}_{i \in M_{src}} \right) = p \left(x_T^j \right) \prod_{t=1}^T p_{\theta} \left(x_t^j | x_{t-1}^j, \Psi \left(\{ e_i^{src} \}_{i \in M_{src}} \right) \right)$$

where Ψ represents the conditional embedding fusion function that integrates information from all source modalities. This formulation allows for flexible conditioning, where any combination of modalities can guide the generation process.

A key innovation in our algorithm is the cross-modal mutual guidance mechanism that facilitates information exchange between concurrently generated modalities. During the progressive denoising process, intermediate representations of each target modality inform the generation of other target modalities, creating a synergistic effect that enhances overall coherence. At each timestep t , the cross-modal guidance is implemented as:

$$\hat{x}_t^j = x_t^j + \sum_{k \in M_{tgt}, k \neq j} \lambda_{j,k} (t) \cdot \Gamma_{j,k} \left(x_t^k, x_t^j \right)$$

where $\lambda_{j,k} (t)$ is a time-dependent weighting function that controls the influence of modality k on modality j , and $\Gamma_{j,k}$ is a cross-modal alignment function that maps information from modality k to modality j . This mutual guidance mechanism enables the model to maintain semantic consistency across simultaneously generated modalities.

To optimize cross-modal consistency, we introduce a novel consistency regularization approach that explicitly enforces semantic alignment between generated modalities. During training, the model minimizes a consistency loss function:

$$\mathcal{L}_{consist} = \sum_{j,k \in M_{tgt}, j \neq k} d_{j,k} \left(\Phi_j \left(x_0^j \right), \Phi_k \left(x_0^k \right) \right)$$

where Φ_j and Φ_k are modal-specific projection functions that map the generated content to a shared semantic space, and $d_{j,k}$ measures the semantic distance between the projected representations. During inference, this consistency is maintained through an iterative refinement process that adjusts the generated content to maximize cross-modal alignment.

Step	Algorithm operation	Input/Output description
1	Modal condition identification	Input: SOURCE modality data $\{ x_i^{src} \}$; Output: Source modality set M_{src} and target modality set M_{tgt}
2	Feature extraction	Input: Source modality data $\{ x_i^{src} \}$; Output: Modality-specific embeddings $\{ e_i^{src} \}$ through corresponding encoders
3	Conditional embedding fusion	Input: Source modality embeddings $\{ e_i^{src} \}$; Output: Unified conditional representation $c_{unified}$ via hierarchical fusion
4	Generation planning	Input: M_{src} , M_{tgt} , generation parameters; Output: Optimal generation sequence π and modality-specific sampling parameters $\{ \theta_j \}$
5	Latent initialization	Input: Target modality specifications; Output: Random noise tensors $\{ x_T^j \}$ for each target modality $j \in M_{tgt}$
6	Progressive denoising	Input: Noise tensors $\{ x_T^j \}$, conditional embedding $c_{unified}$; Output: Intermediate representations $\{ x_t^j \}$ at each timestep
7	Cross-modal consistency enforcement	Input: Intermediate representations $\{ x_t^j \}$; Output: Consistency-adjusted representations $\{ \hat{x}_t^j \}$ with enhanced semantic alignment
8	Adaptive refinement	Input: Consistency-adjusted representations $\{ \hat{x}_t^j \}$, quality assessment metrics; Output: Refined representations with quality-focused adjustments
9	Final decoding	Input: Denoised latent representations $\{ x_0^j \}$; Output: Generated content in target modalities $\{ \hat{y}_j \}$
10	Modal coherence validation	Input: Generated outputs $\{ \hat{y}_j \}$, source inputs $\{ x_i^{src} \}$; Output: Coherence scores and potential regeneration flags for low-quality outputs

Table 3. Multi-modal collaborative generation algorithm flow.

The algorithm implements an adaptive sampling strategy that dynamically adjusts the denoising process based on generation quality and consistency metrics. The sampling step size for modality j at timestep t is computed as:

$$\alpha_t^j = \alpha_{base} \cdot (1 + \beta_q \cdot Q_t^j + \beta_c \cdot C_t^j)$$

where α_{base} is the base step size, Q_t^j is a quality assessment function for modality j , C_t^j is a consistency evaluation function, and β_q and β_c are weighting parameters that balance the influence of quality and consistency on the sampling process.

The generation planning module (Step 4) determines the optimal generation sequence π based on the information density and specificity of the available source modalities. For example, when text and image modalities are provided as conditions, the algorithm may first generate the audio modality from both text and image, leveraging the rich semantic information from both sources. This dynamic planning approach enables the algorithm to adapt to various conditional generation scenarios while maximizing the quality and coherence of the generated content.

Through the integration of flexible conditional generation, cross-modal mutual guidance, consistency optimization, and adaptive sampling strategies, the proposed MCG algorithm achieves robust performance across diverse generation scenarios, effectively addressing the challenges of maintaining semantic consistency and high generation quality across the three modalities.

Cross-domain application experiments and evaluation

Experimental design and evaluation metrics

To comprehensively evaluate the proposed multi-modal diffusion model’s performance in cross-domain applications, we establish a rigorous experimental framework comprising carefully selected datasets, standardized preprocessing methodologies, and multi-faceted evaluation metrics. All experiments were conducted on a high-performance computing cluster equipped with 8×NVIDIA A100 (80GB) GPUs, 512GB RAM, and 64-core AMD EPYC 7763 processors. The implementation was developed using PyTorch 2.0.1 with CUDA 11.8, and all models were trained using mixed precision to optimize computational efficiency. The full model training required approximately 86 h on our 8-GPU cluster, with a batch size of 64 per GPU. Average inference times for different generation pathways are: Text→Image+Audio (4.2 s), Image→Text+Audio (3.8 s), and Audio→Text+Image (5.1 s). Memory requirements are 78GB during training (distributed across GPUs) and 12GB for inference. Compared to baseline methods, our approach requires approximately 15% more computation time but demonstrates 11.65% higher cross-modal coherence, representing a favorable performance-efficiency tradeoff.

The experimental datasets were selected to ensure comprehensive coverage of diverse domains and modality interactions, as detailed in Table 4. These datasets collectively provide a robust foundation for evaluating the model’s generalization capabilities across various application scenarios and content types.

For comprehensive evaluation, we developed a multi-faceted metric system that assesses both modality-specific quality and cross-modal semantic consistency. The modality-specific metrics include:

For image generation: Fréchet Inception Distance (FID), Inception Score (IS), CLIP Score (CS), and Kernel Inception Distance (KID). The FID metric quantifies the statistical distance between generated and real image distributions:

$$FID = \| \mu_r - \mu_g \|^2 + \text{Tr} \left(\Sigma_r + \Sigma_g - 2\sqrt{\Sigma_r \Sigma_g} \right)$$

where μ_r, μ_g and Σ_r, Σ_g are the means and covariance matrices of the real and generated image feature distributions, respectively. To ensure statistical reliability, we calculate FID using 10,000 test samples rather than the previously mentioned 1,000. We also report KID, which is less sensitive to sample size, and LPIPS for perceptual similarity assessment.

Dataset name	modality type	Sample count	Feature dimensions	Preprocessing method
MS-COCO-audio	Text-image-audio	123,287	Text: 77×768, Image: 3×256×256, Audio: 1×16,000×1	Text: BERT tokenization with max length 77; Image: Center-crop, resize to 256×256, normalize to [−1, 1]; Audio: 16 kHz mono resampling, 10-s segments
AudioCaps	Text-audio	49,838	Text: 77×768, Audio: 1×16,000×1	Text: BERT tokenization with max length 77; Audio: 16 kHz mono resampling, 10-s segments with random cropping
Flickr8k-audio	Text-image-audio	8000	Text: 77×768, Image: 3×256×256, Audio: 1×16,000×1	Text: BERT tokenization with max length 77; Image: Center-crop, resize to 256×256, normalize to [−1, 1]; Audio: 16 kHz mono resampling, 10-s segments
HowTo100M-subset	Text-image-audio	50,000	Text: 77×768, Image: 3×256×256, Audio: 1×16,000×1	Text: Transcript cleaning and BERT tokenization; Image: Frame extraction at 1 fps, resize to 256×256; Audio: 16 kHz mono resampling with environmental noise filtering
MusicCaps	Text-audio	5521	Text: 77×768, Audio: 1×16,000×1	Text: BERT tokenization with max length 77; Audio: 16 kHz mono resampling, 10-s segments with volume normalization

Table 4. Experimental datasets and preprocessing methods.

For text generation: BLEU, ROUGE-L, and BERTScore, measuring lexical overlap and semantic similarity with reference texts.
For audio generation: Fréchet Audio Distance (FAD), Kullback–Leibler Divergence (KLD), and Mean Opinion Score (MOS) from human evaluations.

To evaluate cross-modal consistency, we introduce a novel Tri-modal Semantic Alignment Score (TSAS) that quantifies the semantic coherence across the three modalities:

$$TSAS = \frac{1}{3} (\text{sim}(e_{\text{text}}, e_{\text{image}}) + \text{sim}(e_{\text{text}}, e_{\text{audio}}) + \text{sim}(e_{\text{image}}, e_{\text{audio}}))$$

where e_{modality} represents the embedding of each modality in the shared semantic space, and $\text{sim}(\cdot, \cdot)$ measures the cosine similarity between these embeddings.

To ensure fair comparison across different methods with potentially different embedding spaces, we employ a two-step normalization process: (1) all embeddings are projected to a common CLIP⁶³ space using linear probes trained on a held-out dataset, and (2) embedding magnitudes are normalized through L2 normalization. This approach addresses the “modality gap” issue highlighted in⁶⁴, ensuring that measured distances reflect semantic relationships rather than embedding space characteristics. We supplement these metrics with human evaluations on 200 samples where annotators rated cross-modal coherence on a 1–5 scale.

Additionally, we measure the Conditional Generation Quality (CGQ) to assess how well the generated content adheres to the conditions provided by source modalities:

$$CGQ_j = \frac{1}{|M_{src}|} \sum_{i \in M_{src}} \text{sim}(e_i^{src}, e_j^{gen})$$

where j represents the target modality, M_{src} is the set of source modalities, and $\text{sim}(e_i^{src}, e_j^{gen})$ measures the semantic similarity between source and generated content.

For baseline comparison, we benchmark our approach against state-of-the-art methods in each generation pathway, including DALL-E 3, AudioLDM, ImageBind, CLIP-VQGAN, AudioCLIP, and Make-An-Audio. To ensure fair comparison, all models were either retrained on our datasets or, when code was unavailable, evaluated using their official implementations with carefully matched input conditions.

The experimental protocol follows a standardized procedure: (1) for each generation pathway (e.g., text → image + audio, image → text + audio), randomly select 1,000 test samples from each dataset; (2) generate outputs using our model and baseline methods; (3) compute all evaluation metrics using the same post-processing procedures; (4) conduct significance testing (paired t-tests with Bonferroni correction) to verify statistical significance of performance differences.

To ensure reproducibility, we document all hyperparameters, random seeds, and implementation details in the supplementary materials. The code repository, pre-trained models, and processed datasets will be made publicly available upon publication.

Media content creation application results

The proposed multi-modal diffusion model demonstrates significant capabilities in media content creation applications, enabling novel creative workflows that leverage the semantic relationships between textual descriptions, visual elements, and audio components. We evaluated the model’s performance across several key media content creation tasks, including music visualization, sound-to-image generation, automated soundtrack composition, and complete audiovisual content synthesis. Table 5 presents a comprehensive comparison of our approach against state-of-the-art baseline methods across these creative application scenarios.

In music visualization applications, our model achieved a Tri-modal Semantic Alignment Score (TSAS) of 0.782, representing a 20.1% improvement over the strongest baseline combination. The model demonstrated a superior ability to capture the emotional tone, rhythm patterns, and thematic elements of musical pieces, translating them into visually coherent and aesthetically appropriate imagery. This improved performance can be attributed to the hierarchical cross-modal alignment mechanism that effectively bridges the representation gap between audio frequency patterns and visual compositional elements. When visualizing complex musical pieces with varying emotional dynamics, the model successfully generated visual sequences that reflected the temporal progression of the audio, maintaining semantic consistency throughout tonal transitions.

Application scenario	Our method (TSAS)	DALL-E 3 + AudioLDM	ImageBind	AudioCLIP + CLIP-VQGAN	Performance improvement (%)
Music visualization	0.782	0.651	0.603	0.587	+ 20.1
Sound-to-image generation	0.756	N/A	0.683	0.615	+ 10.7
Text-guided soundtrack composition	0.814	0.759	N/A	0.692	+ 7.2
Image-guided audio generation	0.743	N/A	0.679	0.618	+ 9.4
Complete audiovisual synthesis	0.801	0.692	0.665	0.643	+ 15.8
Interactive story illustration	0.775	0.726	N/A	0.681	+ 6.7

Table 5. Media content creation application performance comparison.

Application scenario	User group	Evaluation dimension	Satisfaction Score (1–5)	Improvement suggestions
Multi-sensory learning materials	K-12 educators (n = 42)	Content relevance	4.3	Enhanced domain-specific vocabulary
Multi-sensory learning materials	Students with learning differences (n = 63)	Engagement	4.5	More interactive elements
Visual-to-audio content translation	Visually impaired users (n = 28)	Information preservation	4.1	Improved spatial relationship descriptions
Audio-to-visual content translation	Hearing impaired users (n = 31)	Contextual accuracy	3.9	Better emotional tone representation
Automated educational resource generation	Higher education faculty (n = 37)	Academic accuracy	3.8	More citation capabilities
Concept visualization for STEM education	STEM educators (n = 29)	Scientific accuracy	4.0	Enhanced technical precision
Language learning materials	ESL instructors (n = 35)	Cultural appropriateness	4.2	More diverse cultural representations

Table 6. Assistive technology and education application user evaluation.

For sound-to-image generation tasks, the model exhibited robust performance in translating environmental sounds, human vocalizations, and instrumental audio into contextually relevant imagery. The generated images demonstrated high semantic fidelity to the source audio, accurately representing spatial arrangements, object interactions, and environmental contexts implied by acoustic properties. This capability can be quantified using the Cross-modal Coherence Index (CCI):

$$CCI = \frac{1}{N} \sum_{i=1}^N \frac{\text{sim}(f_A(a_i), f_I(I_i))}{\text{sim}(f_A(a_i), f_I(I_r))}$$

where f_A and f_I are audio and image encoders, a_i is the input audio, I_i is the generated image, and I_r is a random image from the test set. Our model achieved an average CCI of 3.78, significantly outperforming baseline methods.

In automated soundtrack composition, the model demonstrated a notable ability to generate contextually appropriate audio content for given textual descriptions or visual scenes. The Emotional Congruence Score (ECS) for our generated soundtracks reached 0.831, calculated as:

$$ECS = \frac{\sum_{j=1}^M \sum_{k=1}^K w_k \cdot \text{match}(e_j^{gen}, e_j^{ref}, k)}{\sum_{j=1}^M \sum_{k=1}^K w_k}$$

where e_j^{gen} and e_j^{ref} represent the emotional profiles of generated and reference audio, k indexes emotion categories, and w_k are importance weights.

The model's most compelling results emerged in complete audiovisual content synthesis, where it generated semantically coherent combinations of images and sound based on textual prompts. The generated content demonstrated not only high individual quality in each modality but also strong cross-modal coherence, creating immersive experiences where visual and auditory elements mutually reinforced narrative and emotional elements. Professional content creators who evaluated the system (n = 24) rated the outputs as “highly usable for professional applications” (4.2/5.0), noting particular strengths in maintaining thematic consistency across modalities.

Despite these achievements, several limitations were observed in creative applications. First, the model occasionally struggled with highly abstract or culturally specific creative concepts, defaulting to more literal interpretations. Second, in scenarios requiring precise temporal synchronization between visual and audio events, the model sometimes produced slight misalignments that required manual correction. Third, while the model excelled at generating content aligned with conventional aesthetic patterns found in the training data, its ability to produce truly novel artistic styles or experimental creative forms remained limited.

These findings suggest that while the proposed model significantly advances the state-of-the-art in multimodal content creation, further research is needed to enhance temporal precision, cultural adaptability, and creative flexibility for professional media production workflows.

Assistive technology and education application evaluation

The integration of text, image, and audio generation capabilities presents significant opportunities for enhancing accessibility and educational effectiveness through multimodal content. We conducted systematic evaluations of our model's performance in assistive technology and educational applications, focusing on multi-sensory learning materials generation, accessibility enhancement for individuals with disabilities, and automated educational resource construction. Table 6 presents detailed user evaluation results across these application domains.

In the domain of multi-sensory learning materials generation, our model demonstrated significant advantages in creating coherent content that leverages complementary modalities to reinforce learning objectives. When generating explanatory content for complex scientific concepts, the model achieved a Cognitive Reinforcement Index (CRI) of 0.76, calculated as:

$$\text{CRI} = \frac{1}{N} \sum_{i=1}^N (\alpha \cdot \text{con}(t_i, v_i) + \beta \cdot \text{con}(t_i, a_i) + \gamma \cdot \text{con}(v_i, a_i))$$

where $\text{con}(\cdot, \cdot)$ measures conceptual alignment between modalities, t_i , v_i , and a_i represent text, visual, and audio components, and α , β , and γ are modality importance weights that sum to 1. This performance reflects a 32% improvement over current educational content generation systems that typically operate on single or dual modalities.

For accessibility applications, the model's cross-modal translation capabilities demonstrated particular promise for individuals with sensory impairments. Visually impaired users rated the system's ability to convert visual content into descriptive text and representative audio at 4.1/5.0, noting significant improvements in information preservation compared to traditional alt-text generation systems. Similarly, hearing-impaired users reported that the visual representations generated from audio content effectively captured both semantic meaning and emotional tone, achieving a Semantic Retention Score (SRS) of:

$$\text{SRS} = \frac{\sum_{j=1}^M w_j \cdot \text{match}(e_j^{\text{orig}}, e_j^{\text{gen}})}{\sum_{j=1}^M w_j}$$

where e_j^{orig} and e_j^{gen} represent semantic elements in original and generated content, w_j are importance weights, and $\text{match}(\cdot, \cdot)$ assesses semantic preservation.

In educational resource construction, the model demonstrated versatility in generating comprehensive materials from minimal inputs. Educators reported that the ability to automatically expand a core concept into coordinated textual explanations, visual illustrations, and audio narrations significantly reduced content development time (average reduction: 63%). The generated materials showed particular strength in maintaining conceptual consistency across modalities, with 87% of educators reporting that visual and audio elements effectively reinforced the textual content rather than presenting contradictory information.

STEM educators highlighted the model's effectiveness in visualizing abstract concepts, with 79% of participants indicating that the generated visual-audio explanations improved student comprehension of complex topics compared to text-only or static image materials. Language learning applications also showed promising results, with ESL instructors noting that the coordinated presentation of pronunciation audio, contextual imagery, and written text provided a more immersive learning experience.

Despite these positive outcomes, several areas for improvement emerged from user feedback. Educational domain experts noted occasional inaccuracies in highly specialized subject areas, suggesting the need for domain-specific fine-tuning. Users with disabilities identified specific accessibility enhancements needed for various impairment profiles. Additionally, cross-cultural applications revealed limitations in representing diverse cultural contexts, particularly for educational content targeted at international learners.

These findings demonstrate the model's substantial potential for enhancing accessibility and educational effectiveness through coordinated multimodal content generation, while also identifying specific pathways for future refinement to address the diverse needs of educational and assistive technology users.

Conclusion

Summary of research contributions

This research advances the state of multimodal content generation through the development of a novel framework for text-image-audio collaborative generation based on diffusion models. The primary technical contributions include: (1) a hierarchical cross-modal alignment network (HiCAN) that establishes unified representations across three distinct modalities while preserving modality-specific characteristics; (2) a cross-modal conditional diffusion model that enables flexible generation pathways through innovative conditional embedding and attention-guided diffusion mechanisms; and (3) a comprehensive multimodal collaborative generation algorithm that implements cross-modal mutual guidance and consistency optimization to ensure semantic coherence across generated modalities.

Experimental evaluations across diverse application domains demonstrate that our approach significantly outperforms existing methods in both generation quality and cross-modal consistency. In media content creation applications, the proposed model achieved an average improvement of 11.65% in tri-modal semantic alignment compared to state-of-the-art baselines, with particularly strong performance in music visualization (+20.1%) and complete audiovisual synthesis (+15.8%). For assistive technology and educational applications, user evaluations confirmed the model's effectiveness in generating high-quality, semantically consistent multimodal content that enhances information accessibility and learning experiences across diverse user groups.

Advantages and limitations

The primary advantages of our approach stem from three key innovations. First, the unified representation learning framework successfully bridges the semantic gap between heterogeneous modalities, enabling more effective information transfer across the text, image, and audio domains. Second, the attention-guided diffusion process provides fine-grained control over cross-modal influences during generation, allowing for more nuanced semantic consistency. Third, the adaptive sampling strategy enhances generation quality by dynamically allocating computational resources based on real-time quality assessment.

Despite these advances, several limitations warrant acknowledgment. First, computational efficiency remains a challenge, with the full tri-modal generation process requiring significant GPU resources that may limit

deployment in resource-constrained environments. Second, the model exhibits varying performance across different generation pathways, with audio-to-image generation demonstrating lower fidelity compared to text-to-image or text-to-audio pathways. Third, while the model performs well on content types represented in the training data, its generalization capabilities for highly specialized domains or culturally specific content require further enhancement.

Additionally, temporal synchronization between static (text, images) and dynamic (audio) modalities presents ongoing challenges, particularly for precise alignment of audio events with visual elements in complex scenarios. User feedback from educational and assistive technology applications also highlighted the need for improved domain adaptation capabilities to address specialized vocabulary and technical concepts in professional contexts.

Future research directions

Several promising research directions emerge from this work. First, exploring efficient architecture designs and optimization strategies could reduce computational requirements while maintaining generation quality, making multimodal systems more accessible for real-world applications. Second, incorporating more sophisticated temporal modeling approaches could enhance synchronization between static and dynamic modalities, particularly valuable for video generation and interactive applications.

Further investigation into few-shot adaptation mechanisms would address the current limitations in domain-specific applications, enabling rapid customization for specialized contexts without extensive retraining. Exploring self-supervised learning approaches could reduce dependence on aligned tri-modal datasets, allowing models to leverage abundant unimodal or bimodal data more effectively. Additionally, developing more nuanced evaluation metrics specifically designed for tri-modal coherence would facilitate more precise assessment of multimodal generation quality.

The integration of user feedback loops and interactive refinement mechanisms represents another promising direction, allowing systems to iteratively improve generated content based on specific user requirements. This could significantly enhance personalization capabilities and address the current limitations in creative flexibility.

From an application perspective, emerging opportunities exist in immersive technologies (VR/AR), healthcare (therapeutic content generation), scientific visualization, and cultural heritage preservation. The framework's ability to generate coordinated multimodal content holds particular promise for creating more inclusive and accessible digital environments that accommodate diverse sensory capabilities and learning preferences.

In conclusion, this research establishes a foundation for tri-modal collaborative generation that significantly advances multimodal content creation capabilities. While challenges remain, the demonstrated performance across diverse application domains confirms the viability of diffusion-based approaches for complex multimodal generation tasks and highlights numerous pathways for continued innovation in this rapidly evolving field.

Data availability

The datasets and code used in this study, including the processed versions of MS-COCO-Audio, AudioCaps, Flickr8k-Audio, HowTo100M-Subset, and MusicCaps, along with the implementation of the proposed models and evaluation metrics, will be made publicly available through a dedicated repository upon publication. Pre-trained model weights and the complete evaluation results are also included in the repository. Additional experimental data generated during the study are available from the corresponding author upon reasonable request.

Received: 7 March 2025; Accepted: 4 June 2025

Published online: 01 July 2025

References

- Ho, J., Jain, A. & Abbeel, P. Denoising diffusion probabilistic models. *Adv. Neural. Inf. Process. Syst.* **33**, 6840–6851 (2020).
- Ramesh, A., Dhariwal, P., Nichol, A., et al. Hierarchical text-conditional image generation with CLIP latents. *arXiv preprint arXiv:2204.06125* (2022).
- Baltrusaitis, T., Ahuja, C. & Morency, L. P. Multimodal machine learning: A survey and taxonomy. *IEEE Trans. Pattern Anal. Mach. Intell.* **41**(2), 423–443 (2019).
- Song, Y., Sohl-Dickstein, J., Kingma, D. P. et al. Score-based generative modeling through stochastic differential equations. In *International Conference on Learning Representations* (Virtual Event, 2021).
- Saharia, C. et al. Photorealistic text-to-image diffusion models with deep language understanding. *Adv. Neural. Inf. Process. Syst.* **35**, 36479–36494 (2022).
- Rombach, R., Blattmann, A., Lorenz, D. et al. High-resolution image synthesis with latent diffusion models. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition* 10684–10695 (IEEE, 2022).
- Liu, Z., Hu, H., Wang, Y. et al. AudioLDM: Text-to-audio generation with latent diffusion models. In *Proceedings of the 40th International Conference on Machine Learning* 21450–21465 (PMLR, 2023).
- Yang, Z. et al. Diffusion models: A comprehensive survey of methods and applications. *ACM Comput. Surv.* **55**(12), 1–41 (2023).
- Huang, X. et al. MultimodalAI: A survey of applications, datasets, and models. *IEEE Trans. Knowl. Data Eng.* **35**(11), 10699–10718 (2023).
- Lu, Y. et al. Multimodal content generation: A survey of methods and applications. *ACM Comput. Surv.* **56**(1), 1–38 (2024).
- Akbari, H., Yuan, L., Qian, R. et al. VATT: Transformers for multimodal self-supervised learning from raw video, audio and text. In *Advances in Neural Information Processing Systems* 24206–24221 (Virtual Conference: NeurIPS, 2021).
- SINGH A, GOSWAMI S, PALANISAMY V, et al. Cross-modal alignment methods in diffusion models: A comparative study[J]. *arXiv preprint arXiv:2304.10685*, 2023.
- Nichol, A., Dhariwal, P. Improved denoising diffusion probabilistic models. In *Proceedings of the 38th International Conference on Machine Learning* 8162–8171 (Virtual Conference: PMLR, 2021).
- Zhou, L. et al. Comprehensive assessment metrics for multimodal content generation. *IEEE Trans. Multimedia* **25**(7), 3745–3760 (2023).
- Li, Y., Hu, H., Li, C. et al. FusionDM: A unified framework for multimodal content creation via hierarchical diffusion processes. In *Proceedings of the IEEE/CVF International Conference on Computer Vision* 16585–16596 (IEEE, 2023).

16. Yan, W., Vemprala, S., Bonatti, R. et al. Discrete Contrastive diffusion for cross-modal music and image generation. In *International Conference on Learning Representations (ICLR)*, 2023).
17. Zhang, D., Li, C., Chen, J. et al. Cross-modal contrastive learning for multimodal diffusion models. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition* 22823–22834 (IEEE, 2023).
18. Sohl-Dickstein, J., Weiss, E., Maheswaranathan, N. et al. Deep unsupervised learning using nonequilibrium thermodynamics. In *Proceedings of the 32nd International Conference on Machine Learning* 2256–2265 (PMLR, 2015).
19. Dhariwal, P. & Nichol, A. Diffusion models beat GANs on image synthesis. *Adv. Neural Inf. Process. Syst.* **34**, 8780–8794 (2021).
20. Kong, L. et al. A theoretical analysis of diffusion model parameter estimation. *J. Mach. Learn. Res.* **24**(169), 1–49 (2023).
21. Yang, Y. et al. Diffusion models: Theory and application to image processing and generation. *IEEE Signal Process. Mag.* **40**(1), 8–27 (2023).
22. Kingma, D. P. et al. Variational diffusion models. *Adv. Neural Inf. Process. Syst.* **34**, 21696–21707 (2021).
23. Liu, H., Shen, Z., Zhu, Y. et al. Noise estimation for diffusion models: Theory and applications. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition* 7854–7864 (IEEE, 2023).
24. Hoogeboom, E. et al. Argmax flows and multinomial diffusion: Learning categorical distributions. *Adv. Neural Inf. Process. Syst.* **34**, 12454–12467 (2021).
25. Xie, M. et al. GAN-based image generation: A comparative analysis with diffusion models. *IEEE Trans. Image Process.* **32**, 4109–4123 (2023).
26. Croitoru, F. et al. Diffusion models in vision: A survey. *IEEE Trans. Pattern Anal. Mach. Intell.* **45**(12), 14584–14604 (2023).
27. Bond-Taylor, S. et al. Deep generative modelling: A comparative review of VAEs, GANs, normalizing flows, energy-based and autoregressive models. *IEEE Trans. Pattern Anal. Mach. Intell.* **45**(4), 4462–4487 (2023).
28. Kingma, D. P. & Welling, M. An introduction to variational autoencoders. *Found. Trends Mach. Learn.* **12**(4), 307–392 (2019).
29. Lam, M. W. et al. Efficient diffusion modeling for audio synthesis. *IEEE/ACM Trans. Audio Speech Lang. Process.* **31**, 3056–3071 (2023).
30. Guo, H. et al. A systematic review on multimodal learning. *IEEE Trans. Neural Netw. Learn. Syst.* **34**(9), 6032–6049 (2023).
31. Liang, P. P. et al. Foundations and trends in multimodal machine learning: Principles, challenges, and open questions. *Proc. IEEE* **111**(9), 913–942 (2023).
32. Zhang, D. et al. Cross-modal translation methods in diffusion models: A survey. *J. Artif. Intell. Res.* **78**, 1189–1231 (2023).
33. Li, G. et al. Toward multimodal image-to-image translation. *IEEE Trans. Neural Netw. Learn. Syst.* **32**(4), 1713–1727 (2020).
34. Deng, C. et al. Multimodal transformer networks for vision-language translation. *IEEE Trans. Pattern Anal. Mach. Intell.* **44**(5), 2555–2570 (2022).
35. Wu, S., Li, Z., Chen, J. et al. AlignDiff: Aligning diverse modalities in diffusion models for multimodal generation. In *Proceedings of the IEEE/CVF International Conference on Computer Vision* 18394–18404 (IEEE, 2023).
36. Chen, T., Kornblith, S., Norouzi, M. et al. A simple framework for contrastive learning of visual representations. In *Proceedings of the 37th International Conference on Machine Learning* 1597–1607 (Virtual Conference: PMLR, 2020).
37. Yang, J. et al. GeneCIS: A benchmark for general conditional image similarity. *Adv. Neural Inf. Process. Syst.* **36**, 26487–26499 (2023).
38. Vaswani, A., Shazeer, N., Parmar, N. et al. Attention is all you need. In *Advances in Neural Information Processing Systems* 5998–6008 (NeurIPS, 2017).
39. Liu, H. et al. Multi-modal attention networks for cross-modal understanding in diffusion models. *IEEE Trans. Multimedia* **25**(8), 4301–4312 (2023).
40. Li, L., Zhou, A., Chen, J. et al. Cross-attention transformers for multimodal semantic correspondence. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition* 6533–6542 (IEEE, 2023).
41. Wu, L., Li, X., Jiang, L. et al. MultiDiffusion: Fusing diffusion paths for controlled image generation. In *Proceedings of the 40th International Conference on Machine Learning* 37868–37882 (PMLR, 2023).
42. Yang, P., Chen, Z., Liu, Z. et al. Mitigating modality collapse in multimodal conditional diffusion models. In *Advances in Neural Information Processing Systems* 28733–28747 (NeurIPS, 2023).
43. Chen, Y. et al. Temporal coherence in video diffusion models. *IEEE Trans. Image Process.* **33**, 1043–1057 (2024).
44. Gafni, O. et al. Make-A-Scene: Scene-based text-to-image generation with human priors. *ACM Trans. Graphics* **42**(1), 1–22 (2023).
45. Ren, S. et al. Universal cross-modal representation disentanglement for image-text semantic alignment. *IEEE Trans. Pattern Anal. Mach. Intell.* **45**(4), 4714–4728 (2023).
46. Yang, W., Yuan, L., Choi, J. et al. MM-Diffusion: Learning multi-modal diffusion models for joint audio and video generation. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition* 22621–22631 (IEEE, 2023).
47. Huang, J., Zhang, D., Kong, L. et al. TriModalDiff: A unified diffusion model for tri-modal generation. *arXiv preprint arXiv:2304.07585* (2023).
48. Wang, X., Zhan, Z., Li, J. et al. SyncDreamer: Generating multimodal 3D objects using synchronized audio-visual diffusion models. In *Advances in Neural Information Processing Systems* 24661–24674 (NeurIPS, 2023).
49. Kim, S. et al. CascadeDM: Cascaded diffusion models for high fidelity cross-modal generation. *J. Mach. Learn. Res.* **24**(235), 1–35 (2023).
50. Zhou, X., Chai, M., Qi, J. et al. Multi-stage diffusion models for cross-modal audio-visual generation. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition* 9538–9548 (IEEE, 2023).
51. Kim, D. et al. Error propagation analysis in generative cascaded pipelines. *IEEE Trans. Pattern Anal. Mach. Intell.* **45**(7), 8451–8465 (2023).
52. Kim, D., Lin, K., Nguyen, A. et al. Joint audio-text model for text-to-audio generation. In *Proceedings of the IEEE/CVF Winter Conference on Applications of Computer Vision* 1083–1092 (IEEE, 2023).
53. Chen, L. et al. M3AE: Multimodal representation learning for image, text, and audio with masked autoencoders. *IEEE Trans. Multimedia* **25**(9), 4989–5003 (2023).
54. Zhang, Z., Zhang, R., Zhao, R. et al. Semantic consistency models for multi-modal content generation. In *Advances in Neural Information Processing Systems* 15672–15685 (NeurIPS, 2023).
55. Huang, X., Xu, J., Lin, K. et al. AudioToken: Adaptation of text-conditioned diffusion models for audio-to-image generation. *arXiv preprint arXiv:2305.13050* (2023).
56. Ranasinghe, K., Naseer, M., Hayat, M. et al. Perceptual gradients for cross-modal understanding in diffusion models. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition* 21632–21642 (IEEE, 2023).
57. Wu, W. et al. Evaluation metrics for multimodal content generation: A comprehensive review. *IEEE Trans. Knowl. Data Eng.* **35**(10), 9876–9892 (2023).
58. Liu, X., Liu, S., Feng, R. et al. Evaluating cross-modal semantic alignment in diffusion models for multimodal content generation. *arXiv preprint arXiv:2310.12962* (2023).
59. Dong, J., Celik, S., Yu, K. et al. Mitigating modality bias in multimodal learning with audio quality adversarial training. In *International Conference on Acoustics, Speech and Signal Processing* 1–5 (IEEE, 2023).
60. Jiang, L., Hao, Y., Wu, L. et al. AudioViewer: Learning audio-visual correspondence for synchronized audio-visual generation. *arXiv preprint arXiv:2306.17160* (2023).
61. Wu, Q., Zhao, T., Chen, Z. et al. MSP: Multi-stage prompting for making pre-trained language models better translators. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics* 14219–14232 (ACL, 2023).

62. Huang, X., Wu, X., Liu, B. et al. HiCo: Hierarchical co-speech gesture synthesis with controllable diffusion models. In *Advances in Neural Information Processing Systems* 7392–7405 (NeurIPS, 2023).
63. Radford, A., Kim, J. W., Hallacy, C. et al. Learning transferable visual models from natural language supervision. In *Proceedings of the 38th International Conference on Machine Learning* 8748–8763 (Virtual Conference: PMLR, 2021).
64. Liang, P., Chen, Y., Huang, Z. et al. Mind the Gap: Understanding the modality gap in multi-modal contrastive representation learning. In *Advances in Neural Information Processing Systems* 12941–12954 (NeurIPS, 2023).

Acknowledgements

This work was supported by the fourth batch of scientific research and innovation team projects of South China Business College Guangdong University of Foreign Studies (Visual Media Intelligent Analysis And Content Security Research Team), Guangdong Provincial Philosophy and Social Science Planning Project (No. GD23XXL10), 2022 Annual Planning Project of China Private Education Association (No.CANFZG22322), Special project in key fields of colleges and universities in Guangdong province (No. 2024ZDZX1041)

Author contributions

JW conceptualized the research, developed the tri-modal joint representation learning framework, implemented the experimental code, and wrote the initial draft. OZ designed the cross-modal conditional diffusion model architecture, supervised the research, analyzed the experimental results, and revised the manuscript. YJ conducted the user evaluation studies in educational and assistive technology applications, analyzed user feedback data, and contributed to the writing of the application evaluation sections. All authors participated in the discussion of results, reviewed the manuscript, and approved the final version.

Declarations

Competing interests

The authors declare that they have no conflict of interest.

Additional information

Correspondence and requests for materials should be addressed to O.Z.

Reprints and permissions information is available at www.nature.com/reprints.

Publisher's note Springer Nature remains neutral with regard to jurisdictional claims in published maps and institutional affiliations.

Open Access This article is licensed under a Creative Commons Attribution-NonCommercial-NoDerivatives 4.0 International License, which permits any non-commercial use, sharing, distribution and reproduction in any medium or format, as long as you give appropriate credit to the original author(s) and the source, provide a link to the Creative Commons licence, and indicate if you modified the licensed material. You do not have permission under this licence to share adapted material derived from this article or parts of it. The images or other third party material in this article are included in the article's Creative Commons licence, unless indicated otherwise in a credit line to the material. If material is not included in the article's Creative Commons licence and your intended use is not permitted by statutory regulation or exceeds the permitted use, you will need to obtain permission directly from the copyright holder. To view a copy of this licence, visit <http://creativecommons.org/licenses/by-nc-nd/4.0/>.

© The Author(s) 2025